
How To Make the Gradients Small Stochastically: Even Faster Convex and Nonconvex SGD*

Zeyuan Allen-Zhu

Microsoft Research AI

Redmond, WA 98052

zeyuan@csail.mit.edu

Abstract

Stochastic gradient descent (SGD) gives an optimal convergence rate when minimizing convex stochastic objectives $f(x)$. However, in terms of making the gradients small, the original SGD does not give an optimal rate, even when $f(x)$ is convex.

If $f(x)$ is convex, to find a point with gradient norm ε , we design an algorithm SGD3 with a near-optimal rate $\tilde{O}(\varepsilon^{-2})$, improving the best known rate $O(\varepsilon^{-8/3})$ of [17]. If $f(x)$ is nonconvex, to find its ε -approximate local minimum, we design an algorithm SGD5 with rate $\tilde{O}(\varepsilon^{-3.5})$, where previously SGD variants only achieve $\tilde{O}(\varepsilon^{-4})$ [6, 14, 30]. This is no slower than the best known stochastic version of Newton’s method in all parameter regimes [27].

1 Introduction

In convex optimization and machine learning, the classical goal is to design algorithms to decrease objective values, that is, to find points x with $f(x) - f(x^*) \leq \varepsilon$. In contrast, the *rate of convergence for the gradients*, that is,

the number of iterations T needed to find a point x with $\|\nabla f(x)\| \leq \varepsilon$,

is a harder problem and sometimes needs new algorithmic ideas [25]. For instance, in the full-gradient setting, accelerated gradient descent alone is suboptimal for this new goal, and one needs additional tricks to get the fastest rate [25]. We review these tricks in Section 1.1.

In the convex (online) *stochastic* optimization, to the best of our knowledge, tight bounds are not yet known for finding points with small gradients. The best recorded rate was $T \propto \varepsilon^{-8/3}$ [17], and it was raised as an open question [1] regarding how to improve it.

In this paper, we design two new algorithms, SGD2 which gives rate $T \propto \varepsilon^{-5/2}$ using Nesterov’s tricks, and SGD3 which gives an even better rate $T \propto \varepsilon^{-2} \log^3 \frac{1}{\varepsilon}$ which is optimal up to log factors.

Motivation. Studying the rate of convergence for the minimizing gradients can be important at least for the following two reasons.

- In many situations, points with small gradients fit better our final goals.

*The full version of this paper can be found on <https://arxiv.org/abs/1801.02982>. When this paper was submitted to NIPS 2018, the “non-convex SGD” results were not included. We encourage the readers to go to our full version to find out these “non-convex SGD” results.

Nesterov [25] considers the dual approach for solving constrained minimization problems. He argued that “the gradient value $\|\nabla f(x)\|$ serves as the measure of feasibility and optimality of the primal solution,” and thus is the better goal for minimization purpose.²

In matrix scaling [7, 10], given a non-negative matrix, one wants to re-scale its rows and columns to make it doubly stochastic. This problem has been applied in image reconstruction, operations research, decision and control, and other scientific disciplines (see survey [20]). The goal for matrix scaling is to find points with small gradients, but not small objectives.

- Designing algorithms to find points with small gradients can help us understand non-convex optimization better and design faster non-convex machine learning algorithms.

Without strong assumptions, non-convex optimization theory is always in terms of finding points with small gradients (i.e., approximate stationary points or local minima). Therefore, to understand non-convex stochastic optimization better, perhaps we should first figure out the best rate for *convex* stochastic optimization. In addition, if new algorithmic ideas are needed, can we also apply them to the non-convex world? We find positive answers to this question, and also obtain better rates for standard non-convex optimization tasks.

1.1 Review: Prior Work on Deterministic Convex Optimization

Suppose $f(x)$ is a Lipschitz smooth convex function with smoothness parameter L . Then, it is well-known that accelerated gradient descent (AGD) [23, 24] finds a point x satisfying $f(x) - f(x^*) \leq \delta$ using $T = O(\frac{\sqrt{L}}{\sqrt{\delta}})$ gradient computations of $\nabla f(x)$. To turn this into a gradient guarantee, we can apply the smoothness property of $f(x)$ which gives $\|\nabla f(x)\|^2 \leq L(f(x) - f(x^*))$. This means

$$\text{AGD converges in rate } T \propto \frac{L}{\varepsilon}.$$

Nesterov [25] proposed two different tricks to improve upon such rate.

Nesterov’s First Trick: GD After AGD. Recall that starting from a point x_0 , if we perform T steps of gradient descent (GD) $x_{t+1} = x_t - \frac{1}{L}\nabla f(x_t)$, then it satisfies $\sum_{t=0}^{T-1} \|\nabla f(x_t)\|^2 \leq L(f(x_0) - f(x^*))$. In addition, if this x_0 is already the output of AGD for another T iterations, then it satisfies $f(x_0) - f(x^*) \leq O(\frac{L}{T^2})$. Putting the two inequalities together, we have $\min_{t=0}^{T-1} \{\|\nabla f(x_t)\|^2\} \leq O(\frac{L^2}{T^3})$. We call this method “GD after AGD,” and

$$\text{“GD after AGD” converges in rate } T \propto \frac{L^{2/3}}{\varepsilon^{2/3}}.$$

Nesterov’s Second Trick: AGD After Regularization. Alternatively, we can also regularize $f(x)$ by defining $g(x) = f(x) + \frac{\sigma}{2}\|x - x_0\|^2$. This new function $g(x)$ is σ -strongly convex, so AGD converges *linearly*, meaning that using $T \propto \frac{\sqrt{L}}{\sqrt{\sigma}} \log \frac{L}{\varepsilon}$ gradients we can find a point x satisfying $\|\nabla g(x)\|^2 \leq L(g(x) - g(x^*)) \leq \varepsilon^2$. If we choose $\sigma \propto \varepsilon$, then this implies $\|\nabla f(x)\| \leq \|\nabla g(x)\| + \varepsilon \leq 2\varepsilon$. We call this method “AGD after regularization,” and

$$\text{“AGD after regularization” converges in rate } T \propto \frac{L^{1/2}}{\varepsilon^{1/2}} \log \frac{L}{\varepsilon}.$$

This is optimal up to a log factor, because first-order methods need $T = \Omega(\sqrt{L/\delta})$ gradient computations to find $f(x) - f(x^*) \leq \delta$ [23], but $f(x) - f(x^*) \leq \|\nabla f(x)\| \cdot \|x - x^*\| \leq O(\|\nabla f(x)\|)$.

1.2 Our Results: Stochastic Convex Optimization

Consider the stochastic setting where the convex objective $f(x) := \mathbb{E}_i[f_i(x)]$ and the algorithm can only compute stochastic gradients $\nabla f_i(x)$ at any point x for a random i . Let T be the number of stochastic gradient computations. It is well-known that stochastic gradient descent (SGD) finds a point x with $f(x) - f(x^*) \leq \delta$ in (see for instance textbooks [8, 18, 26])

$$T = O\left(\frac{\mathcal{V}}{\delta^2}\right) \text{ iterations} \quad \text{or} \quad T = O\left(\frac{\mathcal{V}}{\sigma\delta}\right) \text{ if } f(x) \text{ is } \sigma\text{-strongly convex.}$$

Both rates are asymptotically optimal in terms of decreasing objective, and \mathcal{V} is an absolute bound on the variance of the stochastic gradients. Using the same argument $\|\nabla f(x)\|^2 \leq L(f(x) - f(x^*))$

²Nesterov [25] studied $\min_{y \in Q} \{g(y) : Ay = b\}$ with convex Q and strongly convex $g(y)$. The dual problem is $\min_x \{f(x)\}$ where $f(x) := \min_{y \in Q} \{g(y) + \langle x, b - Ay \rangle\}$. Let $y^*(x) \in Q$ be the (unique) minimizer of the internal problem, then $g(y^*(x)) - f(x) = \langle x, \nabla f(x) \rangle \leq \|x\| \cdot \|\nabla f(x)\|$.

	algorithm	gradient complexity T	2nd-order smooth
online convex	SGD (naive)	$O(\varepsilon^{-4})$ (folklore, see Theorem 4.2)	no
	SGD1 (SGD after SGD)	$O(\varepsilon^{-8/3})$ (see [17] or Theorem 1)	
	SGD2 (SGD after regularization)	$O(\varepsilon^{-5/2})$ (see Theorem 2)	
	SGD3 (SGD + recursive regularization)	$O(\varepsilon^{-2} \cdot \log^3 \frac{1}{\varepsilon})$ (see Theorem 3)	
online strongly convex	SGD ^{sc} (naive)	$O(\varepsilon^{-2} \cdot \kappa)$ (see Theorem 4.2)	
	SGD1 ^{sc} (SGD after SGD)	$O(\varepsilon^{-2} \cdot \kappa^{1/2})$ (see Theorem 1)	
	SGD3 ^{sc} (SGD + recursive regularization)	$O(\varepsilon^{-2} \cdot \log^3 \kappa)$ (see Theorem 3)	
online nonconvex (σ -nonconvex)	SGD (naive)	$O(\varepsilon^{-4})$ (see [16])	
	SCSG	$O(\varepsilon^{-10/3})$ (see [21])	
	SGD4	$O(\varepsilon^{-2} + \sigma \varepsilon^{-4})$ (see Theorem 4)	
	Natasha1.5	$O(\varepsilon^{-3} + \sigma^{1/3} \varepsilon^{-10/3})$ (see [3])	
	SGD variants	$\tilde{O}(\varepsilon^{-4})$ (see [6, 14, 30])	needed
	SGD5	$\tilde{O}(\varepsilon^{-3.5})$ (see Theorem 5)	
	cubic Newton	$\tilde{O}(\varepsilon^{-3.5})$ (see [27])	
	Natasha2	$O(\varepsilon^{-3.25})$ (see [3])	

Table 1: Comparison of first-order *online stochastic* methods for finding $\|\nabla f(x)\| \leq \varepsilon$. Following tradition, in these bounds, we hide variance and smoothness parameters in big- O and only show the dependency on ε , the condition number $\kappa = \frac{L}{\sigma} \geq 1$ (if the objective is σ -strongly convex), or the nonconvexity parameter σ .

as before, SGD finds a point x with $\|\nabla f(x)\| \leq \varepsilon$ in

$$T = O\left(\frac{L^2 \mathcal{V}}{\varepsilon^4}\right) \text{ iterations} \quad \text{or} \quad T = O\left(\frac{L \mathcal{V}}{\sigma \varepsilon^2}\right) \text{ if } f(x) \text{ is } \sigma\text{-strongly convex.} \quad (\text{SGD})$$

These rates are not optimal. We investigate three approaches to improve such rates.

New Approach 1: SGD after SGD. Recall in Nesterov’s first trick, he replaced the use of the inequality $\|\nabla f(x)\|^2 \leq L(f(x) - f(x^*))$ by T steps of gradient descent. In the stochastic setting, can we replace this inequality with T steps of SGD? We call this algorithm SGD1 and prove that

Theorem 1 (informal). *For convex stochastic optimization, SGD1 finds x with $\|\nabla f(x)\| \leq \varepsilon$ in*

$$T = O\left(\frac{L^{2/3} \mathcal{V}}{\varepsilon^{8/3}}\right) \text{ iterations} \quad \text{or} \quad T = O\left(\frac{L^{1/2} \mathcal{V}}{\sigma^{1/2} \varepsilon^2}\right) \text{ if } f(x) \text{ is } \sigma\text{-strongly convex.} \quad (\text{SGD1})$$

The rate $T \propto \varepsilon^{-8/3}$, in the special case of unconstrained minimization, was first discovered by Ghadimi and Lan [17] using a more complicated algorithm. The rate $T \propto \frac{1}{\sigma^{1/2} \varepsilon^2}$ does not seem to be known before.

New Approach 2: SGD after regularization. Recall that in Nesterov’s second trick, he defined $g(x) = f(x) + \frac{\sigma}{2} \|x - x_0\|^2$ as a regularized version of $f(x)$, and applied the strongly-convex version of AGD to minimize $g(x)$. Can we apply this trick to the stochastic setting?

Note the parameter σ has to be on the magnitude of ε because $\nabla g(x) = \nabla f(x) + \sigma(x - x_0)$ and we wish to make sure $\|\nabla f(x)\| = \|\nabla g(x)\| \pm \varepsilon$. Therefore, if we apply SGD1 to minimize $g(x)$ to find a point $\|\nabla g(x)\| \leq \varepsilon$, the convergence rate is $T \propto \frac{1}{\sigma^{1/2} \varepsilon^2} = \frac{1}{\varepsilon^{2.5}}$. We call this algorithm SGD2.

Theorem 2 (informal). *For convex stochastic optimization, SGD2 finds x with $\|\nabla f(x)\| \leq \varepsilon$ in*

$$T = O\left(\frac{L^{1/2} \mathcal{V}}{\varepsilon^{5/2}}\right) \text{ iterations.} \quad (\text{SGD2})$$

Again, this $T \propto \frac{1}{\varepsilon^{5/2}}$ rate does not seem to be known before.

New Approach 3: SGD and recursive regularization. In the second approach above, the $\varepsilon^{0.5}$ sub-optimality gap is due to the choice of $\sigma \propto \varepsilon$ which ensures $\|\sigma(x - x_0)\| \leq \varepsilon$.

Intuitively, if x_0 were sufficiently close to x^* (and thus were also close to the approximate minimizer x), then we could choose $\sigma \gg \varepsilon$ so that $\|\sigma(x - x_0)\| \leq \varepsilon$ still holds. In other words, an appropriate *warm start* x_0 could help us break the $\varepsilon^{-2.5}$ barrier and get a better convergence rate. However, how to find such x_0 ? We find it by constructing a “less warm” starting point and so on. This process is summarized by the following algorithm which recursively finds the warm starts.

Starting from $f^{(0)}(x) := f(x)$, we define $f^{(s)}(x) := f^{(s-1)}(x) + \frac{\sigma_s}{2}\|x - \hat{x}_s\|^2$ where $\sigma_s = 2\sigma_{s-1}$ and \hat{x}_s is an approximate minimizer of $f^{(s-1)}(x)$ that is simply calculated from the naive SGD. We call this method SGD3, and prove that

Theorem 3 (informal). *For convex stochastic optimization, SGD3 finds x with $\|\nabla f(x)\| \leq \varepsilon$ in*

$$T = O\left(\frac{\log^3(L/\varepsilon) \cdot \mathcal{V}}{\varepsilon^2}\right) \text{ iterations} \quad \text{or} \quad T = O\left(\frac{\log^3(L/\sigma) \cdot \mathcal{V}}{\varepsilon^2}\right) \text{ if } f(x) \text{ is } \sigma\text{-strongly convex. (SGD3)}$$

Our new rates in Theorem 3 not only improve the best known result of $T \propto \varepsilon^{-8/3}$, but also are near optimal because $\Omega(\mathcal{V}/\varepsilon^2)$ is clearly a lower bound: even to decide whether a point x has $\|\nabla f(x)\| \leq \varepsilon$ or $\|\nabla f(x)\| > 2\varepsilon$ requires $\Omega(\mathcal{V}/\varepsilon^2)$ samples of the stochastic gradient. Perhaps interestingly, our dependence on the smoothness parameter L (or the condition number $\kappa := L/\sigma$ if strongly convex) is only polylogarithmic, as opposed to polynomial in all previous results.

1.3 Roadmap

We introduce notions in Section 2 and formalize the convex problem in Section 3. We review classical (convex) SGD theorems with objective decrease in Section 4. We give an auxiliary lemma in Section 5 show our SGD3 results in Section 6. We apply our techniques to non-convex optimization and give algorithms SGD4 and SGD5 in Section 7. We discuss more related work in Appendix A, and show our results on SGD1 and SGD2 respectively in Appendix B and Appendix C.

2 Preliminaries

Throughout this paper, we denote by $\|\cdot\|$ the Euclidean norm. We use $i \in_R [n]$ to denote that i is generated from $[n] = \{1, 2, \dots, n\}$ uniformly at random. We denote by $\nabla f(x)$ the gradient of function f if it is differentiable, and $\partial f(x)$ any subgradient if f is only Lipschitz continuous. We denote by $\mathbb{I}[\text{event}]$ the indicator function of probabilistic events.

We denote by $\|\mathbf{A}\|_2$ the spectral norm of matrix \mathbf{A} . For symmetric matrices \mathbf{A} and \mathbf{B} , we write $\mathbf{A} \succeq \mathbf{B}$ to indicate that $\mathbf{A} - \mathbf{B}$ is positive semidefinite (PSD). Therefore, $\mathbf{A} \succeq -\sigma \mathbf{I}$ if and only if all eigenvalues of \mathbf{A} are no less than $-\sigma$. We denote by $\lambda_{\min}(\mathbf{A})$ and $\lambda_{\max}(\mathbf{A})$ the minimum and maximum eigenvalue of a symmetric matrix \mathbf{A} .

Recall some definitions on strong convexity and smoothness (and they have other equivalent definitions, see textbook [23]).

Definition 2.1. *For a function $f: \mathbb{R}^d \rightarrow \mathbb{R}$,*

- *f is σ -strongly convex if $\forall x, y \in \mathbb{R}^d$, it satisfies $f(y) \geq f(x) + \langle \partial f(x), y - x \rangle + \frac{\sigma}{2}\|x - y\|^2$.*
- *f is of σ -bounded nonconvexity (or σ -nonconvex for short) if $\forall x, y \in \mathbb{R}^d$, it satisfies $f(y) \geq f(x) + \langle \partial f(x), y - x \rangle - \frac{\sigma}{2}\|x - y\|^2$.³*
- *f is L -Lipschitz smooth (or L -smooth for short) if $\forall x, y \in \mathbb{R}^d$, $\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$.*
- *f is L_2 -second-order smooth if $\forall x, y \in \mathbb{R}^d$, it satisfies $\|\nabla^2 f(x) - \nabla^2 f(y)\|_2 \leq L_2\|x - y\|$.*

Definition 2.2. *For composite function $F(x) = \psi(x) + f(x)$ where $\psi(x)$ is proper convex, given a parameter $\eta > 0$, the gradient mapping of $F(\cdot)$ at point x is*

$$\mathcal{G}_{F,\eta}(x) := \frac{1}{\eta}(x - x^+) \quad \text{where} \quad x^+ = \arg \min_y \left\{ \psi(y) + \langle \nabla f(x), y \rangle + \frac{1}{2\eta}\|y - x\|^2 \right\}$$

In particular, if $\psi(\cdot) \equiv 0$, then $\mathcal{G}_{F,\eta}(x) \equiv \nabla f(x)$.

Recall the following property about gradient mapping —see for instance [29, Lemma 3.7])

³Previous authors also refer to this notion as “approximate convex”, “almost convex”, “hypo-convex”, “semi-convex”, or “weakly-convex.” We call it σ -nonconvex to stress the point that σ can be as large as L (any L -smooth function is automatically L -nonconvex).

Lemma 2.3. Let $F(x) = \psi(x) + f(x)$ where $\psi(x)$ is proper convex and $f(x)$ is σ -strongly convex and L -smooth. For every $x, y \in \{x \in \mathbb{R}^d : \psi(x) < +\infty\}$, letting $x^+ = x - \eta \cdot \mathcal{G}_{F,\eta}(x)$, we have

$$\forall \eta \in (0, \frac{1}{L}] : \quad F(y) \geq F(x^+) + \langle \mathcal{G}_{F,\eta}(x), y - x \rangle + \frac{\eta}{2} \|\mathcal{G}_{F,\eta}(x)\|^2 + \frac{\sigma}{2} \|y - x\|^2 .$$

The following definition and properties of Fenchel dual for convex functions is classical, and can be found for instance in the textbook [26].

Definition 2.4. Given proper convex function $h(y)$, its Fenchel dual $h^*(\beta) := \max_y \{y^\top \beta - h(y)\}$.

Proposition 2.5. $\nabla h^*(\beta) = \arg \max_y \{y^\top \beta - h(y)\}$.

Proposition 2.6. If $h(\cdot)$ is σ -strongly convex, then $h^*(\cdot)$ is $\frac{1}{\sigma}$ -smooth.

3 Problem Formalization

Throughout this paper (except our nonconvex application Section 7), we minimize *convex* stochastic composite objective:

$$\min_{x \in \mathbb{R}^d} \{F(x) = \psi(x) + f(x) := \psi(x) + \frac{1}{n} \sum_{i \in [n]} f_i(x)\} , \quad (3.1)$$

where

1. $\psi(x)$ is proper convex (a.k.a. the proximal term),
2. $f_i(x)$ is differentiable for every $i \in [n]$,
3. $f(x)$ is L -smooth and σ -strongly convex for some $\sigma \in [0, L]$ that could be zero,
4. n can be very large or even infinite (so $f(x) = \mathbb{E}_i[f_i(x)]$),⁴ and
5. the stochastic gradients $\nabla f_i(x)$ have a bounded variance (over the domain of $\psi(\cdot)$), that is

$$\forall x \in \{y \in \mathbb{R}^d \mid \psi(y) < +\infty\} : \quad \mathbb{E}_{i \in [n]} \|\nabla f(x) - \nabla f_i(x)\|^2 \leq \mathcal{V} .$$

We emphasize that the above assumptions are all classical.

In the rest of the paper, we define T , the gradient complexity, as the number of computations of $\nabla f_i(x)$. We search for points x so that the gradient mapping $\|\mathcal{G}_{F,\eta}(x)\| \leq \varepsilon$ for any $\eta \approx \frac{1}{L}$. Recall from Definition 2.2 that if there is no proximal term (i.e., $\psi(x) \equiv 0$), then $\mathcal{G}_{F,\eta}(x) = \nabla f(x)$ for any $\eta > 0$. We want to study the best tradeoff between the gradient complexity T and the error ε .

We say an algorithm is *online* if its gradient complexity T is independent of n . This tackles the big-data scenario when n is extremely large or even infinite (i.e., $f(x) = \mathbb{E}_i[f_i(x)]$ for some random variable i). The stochastic gradient descent (SGD) method and all of its variants studied in this paper are online. In contrast, GD, AGD [23, 24], and Katyusha [2] are offline methods because their gradient complexity depends on n (see Table 2 in appendix).

4 Review: SGD with Objective Value Convergence

Recall that stochastic gradient descent (SGD) repeatedly performs *proximal updates* of the form

$$x_{t+1} = \arg \min_{y \in \mathbb{R}^d} \{\psi(y) + \frac{1}{2\alpha} \|y - x_t\|^2 + \langle \nabla f_i(x_t), y \rangle\} ,$$

where $\alpha > 0$ is some learning rate, and i is chosen in $1, 2, \dots, n$ uniformly at random per iteration. Note that if $\psi(y) \equiv 0$ then $x_{t+1} = x_t - \alpha \nabla f_i(x_t)$. For completeness' sake, we summarize it in Algorithm 1. If $f(x)$ is also known to be strongly convex, to get the tightest convergence rate, one can repeatedly apply SGD with decreasing learning rate α [19]. We summarize this algorithm as SGD^{sc} in Algorithm 2.

The following theorem describes the rates of convergence in objective values for SGD and SGD^{sc} respectively. Their proofs are classical (and included in Appendix D); however, for our exact statements, we cannot find them recorded anywhere.⁵

⁴All of the results in this paper apply to the case when n is infinite, because we focus on online methods. However, we still introduce n to simplify notations.

⁵In the special case $\psi(x) \equiv 0$, Theorem 4.1(a) and 4.1(b) are folklore (see for instance [26]). If $\psi(x) \not\equiv 0$, Theorem 4.1(a) is recorded when $\psi(x)$ is Lipschitz or smooth [13], but we would not like to impose such assumptions. A variant of Theorem 4.1(b) is recorded for the accelerated version of SGD [15], but with a

Algorithm 1 $\text{SGD}(F, x_0, \alpha, T)$

Input: function $F(x) = \psi(x) + \frac{1}{n} \sum_{i=1}^n f_i(x)$; initial vector x_0 ; learning rate $\alpha > 0$; $T \geq 1$.
 \diamond if $f(x) = \frac{1}{n} \sum_{i=1}^n f_i(x)$ is L -smooth, optimal choice $\alpha = \Theta(\min\{\frac{\|x_0 - x^*\|}{\sqrt{\mathcal{V}T}}, \frac{1}{L}\})$
1: **for** $t = 0$ **to** $T - 1$ **do**
2: $i \leftarrow$ a random index in $[n]$;
3: $x_{t+1} \leftarrow \arg \min_{y \in \mathbb{R}^d} \{\psi(y) + \frac{1}{2\alpha} \|y - x_t\|^2 + \langle \nabla f_i(x_t), y \rangle\}$;
4: **return** $\bar{x} = \frac{x_1 + \dots + x_T}{T}$.

Algorithm 2 $\text{SGD}^{\text{sc}}(F, x_0, \sigma, L, T)$

Input: function $F(x) = \psi(x) + \frac{1}{n} \sum_{i=1}^n f_i(x)$; initial vector x_0 ; parameters $0 < \sigma \leq L$; $T \geq \frac{L}{\sigma}$.
 \diamond $f(x)$ is σ -strongly convex and $f(x) = \frac{1}{n} \sum_{i=1}^n f_i(x)$ is L -smooth
1: **for** $t = 1$ **to** $N = \lfloor \frac{T}{8L/\sigma} \rfloor$ **do** $x_t \leftarrow \text{SGD}(F, x_{t-1}, \frac{1}{2L}, \frac{4L}{\sigma})$;
2: **for** $k = 1$ **to** $K = \lfloor \log_2(\sigma T / 16L) \rfloor$ **do** $x_{N+k} \leftarrow \text{SGD}(F, x_{N+k-1}, \frac{1}{2^k L}, \frac{2^{k+2}L}{\sigma})$;
3: **return** $\bar{x} = x_{N+K}$.

Theorem 4.1. Let $x^* \in \arg \min_x \{F(x)\}$. To solve Problem (3.1) given a starting vector $x_0 \in \mathbb{R}^d$,
(a) $\text{SGD}(F, x_0, \alpha, T)$ outputs \bar{x} satisfying $\mathbb{E}[F(\bar{x})] - F(x^*) \leq \frac{\alpha \mathcal{V}}{2(1-\alpha L)} + \frac{\|x_0 - x^*\|^2}{2\alpha T}$ as long as $\alpha < 1/L$. In particular, if α is tuned optimally, it satisfies

$$\mathbb{E}[F(\bar{x})] - F(x^*) \leq O\left(\frac{L\|x_0 - x^*\|^2}{T} + \frac{\sqrt{\mathcal{V}}\|x_0 - x^*\|}{\sqrt{T}}\right).$$

(b) If $f(x)$ is σ -strongly convex and $T \geq \frac{L}{\sigma}$, then $\text{SGD}^{\text{sc}}(F, x_0, \sigma, L, T)$ outputs \bar{x} satisfying

$$\mathbb{E}[F(\bar{x})] - F(x^*) \leq O\left(\frac{\mathcal{V}}{\sigma T}\right) + \left(1 - \frac{\sigma}{L}\right)^{\Omega(T)} \sigma \|x_0 - x^*\|^2.$$

As a sanity check, if $\mathcal{V} = 0$, the convergence rate of SGD matches that of GD. (However, if $\mathcal{V} = 0$, one can apply accelerated gradient descent of Nesterov [22, 23] instead for a faster rate.)

To turn Theorem 4.1 into a rate of convergence for the gradients, we can simply apply Lemma 2.3 which implies

$$\forall \eta \in (0, \frac{1}{L}]: \quad \frac{\eta}{2} \|\mathcal{G}_{F,\eta}(\bar{x})\|^2 \leq F(\bar{x}) - F(\bar{x}^+) \leq F(\bar{x}) - F(x^*). \quad (4.1)$$

Theorem 4.2. Let $x^* \in \arg \min_x \{F(x)\}$. To solve Problem (3.1) given a starting vector $x_0 \in \mathbb{R}^d$ and any $\eta = \frac{C}{L}$ where $C \in (0, 1]$ is some absolute constant,

(a) SGD outputs \bar{x} satisfying $\mathbb{E}[\|\mathcal{G}_{F,\eta}(\bar{x})\|^2] \leq O\left(\frac{L^2\|x_0 - x^*\|^2}{T} + \frac{L\sqrt{\mathcal{V}}\|x_0 - x^*\|}{\sqrt{T}}\right)$.

(b) if $T \geq \frac{L}{\sigma}$, then SGD^{sc} outputs \bar{x} satisfying $\mathbb{E}[\|\mathcal{G}_{F,\eta}(\bar{x})\|^2] \leq O\left(\frac{L\mathcal{V}}{\sigma T}\right) + \left(1 - \frac{\sigma}{L}\right)^{\Omega(T)} \sigma L \|x_0 - x^*\|^2$.

Corollary 4.3. Hiding $\mathcal{V}, L, \|x_0 - x^*\|$ in the big- O notation, classical SGD finds x with

$$F(x) - F(x^*) \leq O(T^{-1/2}) \quad \|\mathcal{G}_{F,\eta}(x)\| \leq O(T^{-1/4}) \quad \text{for Problem (3.1), or}$$

$$F(x) - F(x^*) \leq O((\sigma T)^{-1}) \quad \|\mathcal{G}_{F,\eta}(x)\| \leq O((\sigma T)^{-1/2}) \quad \text{if } f(\cdot) \text{ is } \sigma\text{-strongly convex for } \sigma > 0.$$

5 An Auxiliary Lemma on Regularization

Consider a regularized objective

$$G(x) := \psi(x) + g(x) := \psi(x) + \left(f(x) + \sum_{s=1}^S \frac{\sigma_s}{2} \|x - \hat{x}_s\|^2\right), \quad (5.1)$$

slightly worse rate $T = O\left(\frac{\mathcal{V}}{\sigma T} + \frac{L\|x_0 - x^*\|^2}{T^2}\right)$. If the readers find either statement explicitly stated somewhere, please let us know and we would love to include appropriate citations.

where $\hat{x}_1, \dots, \hat{x}_S$ are fixed vectors in \mathbb{R}^d . The following lemma says that, if we find an approximate stationary point x of $G(x)$, then it is also an approximate stationary point of $F(x)$ up to some additive error.

Lemma 5.1. *Suppose $\psi(x)$ is proper convex and $f(x)$ is convex and L -smooth. By definition, $g(x)$ is $\tilde{\sigma}$ -strongly convex with $\tilde{\sigma} := \sum_{s=1}^S \sigma_s$. Let x^* be the unique minimizer of $G(y)$ in (5.1), and x be an arbitrary vector in the domain of $\{x \in \mathbb{R}^d : \psi(x) < +\infty\}$. Then, for every $\eta \in (0, \frac{1}{L+\tilde{\sigma}}]$, we have*

$$\|\mathcal{G}_{F,\eta}(x)\| \leq \sum_{s=1}^S \sigma_s \|x^* - \hat{x}_s\| + 3\|\mathcal{G}_{G,\eta}(x)\|.$$

Remark 5.2. Lemma 5.1 should be easy to prove in the special case of $\psi(x) \equiv 0$. Indeed,

$$\begin{aligned} \|\nabla f(x)\| &= \|\nabla g(x) + \sum_s \sigma_s (x - \hat{x}_s)\| \stackrel{\textcircled{1}}{\leq} \|\nabla g(x)\| + \sum_s \sigma_s \|x - \hat{x}_s\| \\ &\stackrel{\textcircled{2}}{\leq} \|\nabla g(x)\| + \sum_s \sigma_s \|x^* - \hat{x}_s\| + \tilde{\sigma} \|x^* - x\| \stackrel{\textcircled{3}}{\leq} 2\|\nabla g(x)\| + \sum_s \sigma_s \|x^* - \hat{x}_s\|. \end{aligned}$$

Above, inequalities $\textcircled{1}$ and $\textcircled{2}$ both use the triangle inequality; and inequality $\textcircled{3}$ is due to the $\tilde{\sigma}$ -strong convexity of $g(x)$ (see for instance [23, Sec. 2.1.3]).

Proof of Lemma 5.1. See full version. □

6 Approach 3: SGD and Recursive Regularization

In this section, add a logarithmic number of regularizers to the objective, each centered at a different but carefully chosen point. Specifically, given parameters $\sigma_1, \dots, \sigma_S > 0$, we define functions

$$F^{(0)}(x) := F(x) \quad \text{and} \quad F^{(s)}(x) := F^{(s-1)}(x) + \frac{\sigma_s}{2} \|x - \hat{x}_s\|^2 \quad \text{for } s = 1, 2, \dots, S$$

where each \hat{x}_s (for $s \geq 1$) is an approximate minimizer of $F^{(s-1)}(x)$.

If $f(x)$ is σ -strongly convex, then we choose $S \approx \log_2 \frac{L}{\sigma}$ and let $\sigma_0 = \sigma$ and $\sigma_s = 2\sigma_{s-1}$. To calculate each \hat{x}_s , we apply SGD^{sc} for $\frac{T}{S}$ iterations. This totals to a gradient complexity of T . We summarize this method as SGD3^{sc} in Algorithm 3.

If $f(x)$ is not strongly convex, then we regularize it by $G(x) = F(x) + \frac{\sigma}{2} \|x - x_0\|^2$ for some small parameter $\sigma > 0$, and then apply SGD3^{sc}. We summarize this final method as SGD3 in Algorithm 4.

We prove the following main theorem:

Theorem 3 (SGD3). *Let $x^* \in \arg \min_x \{F(x)\}$. To solve Problem (3.1) given a starting vector $x_0 \in \mathbb{R}^d$ and any $\eta = \frac{C}{L}$ for some absolute constant $C \in (0, 1]$.*

- (a) *If $f(x)$ is σ -strongly convex for $\sigma \in (0, L]$ and $T \geq \frac{L}{\sigma} \log \frac{L}{\sigma}$, then SGD3^{sc}(F, x_0, σ, L, T) outputs \bar{x} satisfying*

$$\mathbb{E}[\|\mathcal{G}_{F,\eta}(\bar{x})\|] \leq O\left(\frac{\sqrt{V} \cdot \log^{3/2} \frac{L}{\sigma}}{\sqrt{T}}\right) + \left(1 - \frac{\sigma}{L}\right)^{\Omega(T/\log(L/\sigma))} \sigma \|x_0 - x^*\|.$$

- (b) *If $\sigma \in (0, L]$ and $T \geq \frac{L}{\sigma} \log \frac{L}{\sigma}$, then SGD3(F, x_0, σ, L, T) outputs \bar{x} satisfying*

$$\mathbb{E}[\|\mathcal{G}_{F,\eta}(\bar{x})\|] \leq O\left(\sigma \|x_0 - x^*\| + \frac{\sqrt{V} \cdot \log^{3/2} \frac{L}{\sigma}}{\sqrt{T}}\right) + \left(1 - \frac{\sigma}{L}\right)^{\Omega(T/\log(L/\sigma))} \sigma \|x_0 - x^*\|.$$

If σ is appropriately chosen, then we find \bar{x} with $\mathbb{E}[\|\mathcal{G}_{F,\eta}(\bar{x})\|] \leq \varepsilon$ in gradient complexity

$$T \leq O\left(\frac{V \cdot \log^3 \frac{L\|x_0 - x^*\|}{\varepsilon}}{\varepsilon^2} + \frac{L\|x_0 - x^*\|}{\varepsilon} \log \frac{L\|x_0 - x^*\|}{\varepsilon}\right).$$

Remark 6.1. All expected guarantees of the form $\mathbb{E}[\|\mathcal{G}_{F,\eta}(\bar{x})\|^2] \leq \varepsilon^2$ or $\mathbb{E}[\|\mathcal{G}_{F,\eta}(\bar{x})\|] \leq \varepsilon$ throughout this paper can be made into high-confidence bound by repeating the algorithm multiple times, each time estimating the value of $\|\mathcal{G}_{F,\eta}(\bar{x})\|$ using roughly $O(\frac{V}{\varepsilon^2})$ stochastic gradient computations, and finally outputting the point \bar{x} that leads to the smallest value $\|\mathcal{G}_{F,\eta}(\bar{x})\|$.

Algorithm 3 SGD3^{sc}(F, x_0, σ, L, T)

Input: function $F(x) = \psi(x) + \frac{1}{n} \sum_{i=1}^n f_i(x)$; initial vector x_0 ; parameters $0 < \sigma \leq L$; number of iterations $T \geq \Omega\left(\frac{L}{\sigma} \log \frac{L}{\sigma}\right)$. $\diamond f(x) = \frac{1}{n} \sum_{i=1}^n f_i(x)$ is σ -strongly convex and L -smooth

- 1: $F^{(0)}(x) := F(x); \hat{x}_0 \leftarrow x_0; \sigma_0 \leftarrow \sigma$;
- 2: **for** $s = 1$ **to** $S = \lfloor \log_2 \frac{L}{\sigma} \rfloor$ **do**
- 3: $\hat{x}_s \leftarrow \text{SGD}^{\text{sc}}(F^{(s-1)}, \hat{x}_{s-1}, \sigma_{s-1}, 3L, \frac{T}{S})$;
- 4: $\sigma_s \leftarrow 2\sigma_{s-1}$;
- 5: $F^{(s)}(x) := F^{(s-1)}(x) + \frac{\sigma_s}{2} \|x - \hat{x}_s\|^2$;
- 6: **return** $\bar{x} = \hat{x}_S$.

Algorithm 4 SGD3(F, x_0, σ, L, T)

Input: function $F(x) = \psi(x) + \frac{1}{n} \sum_{i=1}^n f_i(x)$; initial vector x_0 ; parameters $L \geq \sigma > 0; T \geq 1$. $\diamond f(x) = \frac{1}{n} \sum_{i=1}^n f_i(x)$ is convex and L -smooth

- 1: $G(x) := F(x) + \frac{\sigma}{2} \|x - x_0\|^2$;
- 2: **return** $\bar{x} \leftarrow \text{SGD3}^{\text{sc}}(G, x_0, \sigma, L + \sigma, T)$.

6.1 Proof of Theorem 3

Before proving Theorem 3, we state a few properties regarding the relationships between the objective-optimality of \hat{x}_s and point distances.

Claim 6.2. Suppose for every $s = 1, \dots, S$ the vector \hat{x}_s satisfies

$$\mathbb{E}[F^{(s-1)}(\hat{x}_s) - F^{(s-1)}(x_{s-1}^*)] \leq \delta_s \quad (6.1)$$

where $x_{s-1}^* \in \arg \min_x \{F^{(s-1)}(x)\}$, then,

- (a) for every $s \geq 1$, $\mathbb{E}[\|\hat{x}_s - x_{s-1}^*\|^2] \leq \mathbb{E}[\|\hat{x}_s - x_{s-1}^*\|^2] \leq \frac{2\delta_s}{\sigma_{s-1}}$,
- (b) for every $s \geq 1$, $\mathbb{E}[\|\hat{x}_s - x_s^*\|^2] \leq \mathbb{E}[\|x_s^* - \hat{x}_s\|^2] \leq \frac{\delta_s}{\sigma_s}$; and
- (c) if $\sigma_s = 2\sigma_{s-1}$ for all $s \geq 1$, then $\mathbb{E}[\sum_{s=1}^S \sigma_s \|x_s^* - \hat{x}_s\|] \leq 4 \sum_{s=1}^S \sqrt{\delta_s \sigma_s}$.

Proof of Claim 6.2.

- (a) $\mathbb{E}[\|\hat{x}_s - x_{s-1}^*\|^2] \stackrel{\textcircled{1}}{\leq} \mathbb{E}[\|\hat{x}_s - x_{s-1}^*\|^2] \stackrel{\textcircled{2}}{\leq} \frac{2}{\sigma_{s-1}} \mathbb{E}[F^{(s-1)}(\hat{x}_s) - F^{(s-1)}(x_{s-1}^*)] \leq \frac{2\delta_s}{\sigma_{s-1}}$. Here, inequality $\textcircled{1}$ is because $\mathbb{E}[X]^2 \leq \mathbb{E}[X^2]$, and inequality $\textcircled{2}$ is due to the strong convexity of $F^{(s-1)}(x)$.

- (b) We derive that

$$\begin{aligned} \sigma_s \|x_s^* - \hat{x}_s\|^2 &\stackrel{\textcircled{1}}{\leq} \frac{\sigma_s}{2} \|x_s^* - \hat{x}_s\|^2 + F^{(s)}(\hat{x}_s) - F^{(s)}(x_s^*) = F^{(s-1)}(\hat{x}_s) - F^{(s-1)}(x_s^*) \\ &\stackrel{\textcircled{2}}{\leq} F^{(s-1)}(\hat{x}_s) - F^{(s-1)}(x_{s-1}^*) . \end{aligned}$$

Here, inequality $\textcircled{1}$ is due to the strong convexity of $F^{(s)}(x)$, and inequality $\textcircled{2}$ is because of the minimality of x_{s-1}^* . Taking expectation we have $\mathbb{E}[\|x_s^* - \hat{x}_s\|^2] \leq \mathbb{E}[\|x_s^* - \hat{x}_s\|^2] \leq \frac{\delta_s}{\sigma_s}$.

- (c) Define $P_t := \sum_{s=1}^t \sigma_s \|x_t^* - \hat{x}_s\|$ for each $t \geq 0, 1, \dots, S$. Then by triangle inequality we have

$$P_s - P_{s-1} \leq \sigma_s \|x_s^* - \hat{x}_s\| + \left(\sum_{t=1}^{s-1} \sigma_t \right) \cdot \|x_s^* - x_{s-1}^*\|$$

Using the parameter choice of $\sigma_s = 2\sigma_{s-1}$, and plugging in Claim 6.2(a) and Claim 6.2(b), we have

$$\mathbb{E}[P_s - P_{s-1}] \leq \sqrt{\delta_s \sigma_s} + \sigma_s \cdot \mathbb{E}[\|x_s^* - \hat{x}_s\| + \|x_{s-1}^* - \hat{x}_s\|] \leq 4\sqrt{\delta_s \sigma_s} . \quad \square$$

Proof of Theorem 3(a). We first note that, when writing $f^{(s-1)}(x) = F^{(s-1)}(x) - \psi(x)$, each $f^{(s-1)}$ is at least σ_{s-1} -strongly convex and $L + \sum_{t=1}^{s-1} \sigma_t \leq 3L$ Lipschitz smooth. Therefore,

applying Theorem 4.1(b), we have

$$\mathbb{E}[F^{(s-1)}(\hat{x}_s) - F^{(s-1)}(x_{s-1}^*)] \leq O\left(\frac{S\mathcal{V}}{\sigma_{s-1}T}\right) + \left(1 - \frac{\sigma_{s-1}}{3L}\right)^{\Omega(T/S)} \mathbb{E}[\sigma_{s-1}\|\hat{x}_{s-1} - x_{s-1}^*\|^2] .$$

If $s = 1$, this means (recalling $\hat{x}_0 = x_0$ and $x_0^* = x^*$)

$$\mathbb{E}[F^{(0)}(\hat{x}_s) - F^{(0)}(x^*)] \leq O\left(\frac{S\mathcal{V}}{\sigma_0T}\right) + \left(1 - \frac{\sigma_0}{L}\right)^{\Omega(T/S)} \sigma_0\|x_0 - x^*\|^2 .$$

If $s > 1$, this means

$$\mathbb{E}[F^{(s-1)}(\hat{x}_s) - F^{(s-1)}(x_{s-1}^*)] \leq O\left(\frac{S\mathcal{V}}{\sigma_{s-1}T}\right) + \left(1 - \frac{\sigma_{s-1}}{L}\right)^{\Omega(T/S)} \mathbb{E}[F^{(s-2)}(\hat{x}_{s-1}) - F^{(s-2)}(x_{s-2}^*)] .$$

Together, this means to satisfy (6.1), it suffices to choose δ_s so that

$$\delta_s = O\left(\frac{S\mathcal{V}}{\sigma_sT}\right) + \left(1 - \frac{\sigma_0}{L}\right)^{\Omega(sT/S)} \sigma_0\|x_0 - x^*\|^2 .$$

Using Lemma 2.3 with $F^{(S-1)}$ and $y = x = \hat{x}_S$, we have $\frac{\eta}{2}\|\mathcal{G}_{F^{(S-1)},\eta}(\hat{x}_S)\|^2 \leq F^{(S-1)}(\hat{x}_S) - F^{(S-1)}(\hat{x}_S^+) \leq F^{(S-1)}(\hat{x}_S) - F^{(S-1)}(x_{S-1}^*)$ and therefore

$$\mathbb{E}[\|\mathcal{G}_{F^{(S-1)},\eta}(\hat{x}_S)\|]^2 \leq \mathbb{E}[\|\mathcal{G}_{F^{(S-1)},\eta}(\hat{x}_S)\|^2] \leq \frac{2\delta_S}{\eta} = O(L\delta_S) .$$

Plugging this into Lemma 5.1 (with $G(x) = F^{(S-1)}(x)$) and Claim 6.2(c), we have

$$\begin{aligned} \mathbb{E}[\|\mathcal{G}_{F,\eta}(\hat{x}_S)\|] &\leq \mathbb{E}\left[\sum_{s=1}^{S-1} \sigma_s\|x_{s-1}^* - \hat{x}_s\| + 3\|\mathcal{G}_{F^{(S-1)},\eta}(\hat{x}_S)\|\right] \leq O\left(\sum_{s=1}^{S-1} \sqrt{\delta_s\sigma_s} + \sqrt{L\delta_S}\right) \\ &= O\left(\sum_{s=1}^S \sqrt{\delta_s\sigma_s}\right) \leq O\left(\frac{S^{3/2}\mathcal{V}^{1/2}}{T^{1/2}}\right) + \left(1 - \frac{\sigma_0}{L}\right)^{\Omega(T/S)} \sigma_0\|x_0 - x^*\| . \quad \square \end{aligned}$$

Proof of Theorem 3(b). Define $G(x) := F(x) + \frac{\sigma}{2}\|x - x_0\|^2$ and let x_G^* be the (unique) minimizer of $G(\cdot)$. Note that x_G^* may be different from x^* which is a minimizer of $F(\cdot)$. Applying Theorem 3(a) on $G(x)$ and Lemma 5.1 with $S = 1$ and $\hat{x}_1 = x_0$, we have

$$\mathbb{E}[\|\mathcal{G}_{F,\eta}(\bar{x})\|] \leq O\left(\sigma\|x_0 - x_G^*\| + \frac{\sqrt{\mathcal{V}} \cdot \log^{3/2} \frac{L}{\sigma}}{\sqrt{T}}\right) + \left(1 - \frac{\sigma}{L}\right)^{\Omega(T/\log(L/\sigma))} \sigma\|x_0 - x_G^*\|$$

Now, by definition $\frac{\sigma}{2}\|x^* - x_0\|^2 - \frac{\sigma}{2}\|x_G^* - x_0\|^2 = (G(x^*) - F(x^*)) + (F(x_G^*) - G(x_G^*)) \geq 0$ so we have $\|x_G^* - x_0\| \leq \|x^* - x_0\|$. This completes the proof. \square

Acknowledgements

We would like to thank Lin Xiao for suggesting reference [29, Lemma 3.7], an anonymous researcher from the Simons Institute for suggesting reference [25], Yurii Nesterov for helpful discussions, Xinyu Weng for discussing the motivations, Sébastien Bubeck, Yuval Peres, and Lin Xiao for discussing notations, Chi Jin for discussing reference [27], and Dmitriy Drusvyatskiy for discussing the notion of Moreau envelope.

References

- [1] Open problem session of “fast iterative methods in optimization” workshop. Simons Institute for the Theory of Computing, UC Berkeley, October 2017.
- [2] Zeyuan Allen-Zhu. Katyusha: The First Direct Acceleration of Stochastic Gradient Methods. In *STOC*, 2017.
- [3] Zeyuan Allen-Zhu. Natasha 2: Faster Non-Convex Optimization Than SGD. In *NIPS*, 2018.
- [4] Zeyuan Allen-Zhu and Elad Hazan. Optimal Black-Box Reductions Between Optimization Objectives. In *NIPS*, 2016.
- [5] Zeyuan Allen-Zhu and Yuanzhi Li. Follow the Compressed Leader: Faster Online Learning of Eigenvectors and Faster MMWU. In *ICML*, 2017.

- [6] Zeyuan Allen-Zhu and Yuanzhi Li. Neon2: Finding Local Minima via First-Order Oracles. In *NIPS*, 2018.
- [7] Zeyuan Allen-Zhu, Yuanzhi Li, Rafael Oliveira, and Avi Wigderson. Much Faster Algorithms for Matrix Scaling. In *FOCS*, 2017.
- [8] Sébastien Bubeck. Convex optimization: Algorithms and complexity. *Foundations and Trends in Machine Learning*, 8(3-4):231–357, 2015.
- [9] Yair Carmon, John C. Duchi, Oliver Hinder, and Aaron Sidford. Accelerated Methods for Non-Convex Optimization. *ArXiv e-prints*, abs/1611.00756, November 2016.
- [10] M. B. Cohen, A. Madry, D. Tsipras, and A. Vladu. Matrix Scaling and Balancing via Box Constrained Newton’s Method and Interior Point Methods. In *FOCS*, pages 902–913, Oct 2017.
- [11] Damek Davis and Dmitriy Drusvyatskiy. Complexity of finding near-stationary points of convex functions stochastically. *ArXiv e-prints*, abs/1802.08556, 2018.
- [12] Damek Davis and Dmitriy Drusvyatskiy. Stochastic subgradient method converges at the rate $o(k^{-1/4})$ on weakly convex functions. *ArXiv e-prints*, abs/1802.02988, 2018.
- [13] John Duchi and Yoram Singer. Efficient Online and Batch Learning Using Forward Backward Splitting. *Journal of Machine Learning Research*, 10:2899–2934, 2009. ISSN 15324435. doi: 10.1561/2400000003.
- [14] Rong Ge, Furong Huang, Chi Jin, and Yang Yuan. Escaping from saddle points—online stochastic gradient for tensor decomposition. In *Proceedings of the 28th Annual Conference on Learning Theory*, COLT 2015, 2015.
- [15] Saeed Ghadimi and Guanghui Lan. Optimal stochastic approximation algorithms for strongly convex stochastic composite optimization I: A generic algorithmic framework. *SIAM Journal on Optimization*, 22(4):1469–1492, 2012.
- [16] Saeed Ghadimi and Guanghui Lan. Stochastic first-and zeroth-order methods for nonconvex stochastic programming. *SIAM Journal on Optimization*, 23(4):2341–2368, 2013.
- [17] Saeed Ghadimi and Guanghui Lan. Accelerated gradient methods for nonconvex nonlinear and stochastic programming. *Mathematical Programming*, pages 1–26, feb 2015. ISSN 0025-5610.
- [18] Elad Hazan. Introduction to online convex optimization. *Foundations and Trends in Optimization*, 2(3-4):157–325, 2016. ISSN 2167-3888. doi: 10.1561/24000000013.
- [19] Elad Hazan and Satyen Kale. Beyond the regret minimization barrier: Optimal algorithms for stochastic strongly-convex optimization. *The Journal of Machine Learning Research*, 15(1): 2489–2512, 2014.
- [20] Martin Idel. A review of matrix scaling and sinkhorn’s normal form for matrices and positive maps. *ArXiv e-prints*, abs/1609.06349, 2016.
- [21] Lihua Lei, Cheng Ju, Jianbo Chen, and Michael I Jordan. Nonconvex Finite-Sum Optimization Via SCSG Methods. In *NIPS*, 2017.
- [22] Yurii Nesterov. A method of solving a convex programming problem with convergence rate $O(1/k^2)$. In *Doklady AN SSSR (translated as Soviet Mathematics Doklady)*, volume 269, pages 543–547, 1983.
- [23] Yurii Nesterov. *Introductory Lectures on Convex Programming Volume: A Basic course*, volume I. Kluwer Academic Publishers, 2004. ISBN 1402075537.
- [24] Yurii Nesterov. Smooth minimization of non-smooth functions. *Mathematical Programming*, 103(1):127–152, December 2005. ISSN 0025-5610. doi: 10.1007/s10107-004-0552-5.
- [25] Yurii Nesterov. How to make the gradients small. *Optima*, 88:10–11, 2012.

- [26] Shai Shalev-Shwartz. Online Learning and Online Convex Optimization. *Foundations and Trends in Machine Learning*, 4(2):107–194, 2012. ISSN 1935-8237.
- [27] Nilesh Tripuraneni, Mitchell Stern, Chi Jin, Jeffrey Regier, and Michael I Jordan. Stochastic Cubic Regularization for Fast Nonconvex Optimization. *ArXiv e-prints*, abs/1711.02838, November 2017.
- [28] Blake Woodworth and Nati Srebro. Tight Complexity Bounds for Optimizing Composite Objectives. In *NIPS*, 2016.
- [29] Lin Xiao and Tong Zhang. A Proximal Stochastic Gradient Method with Progressive Variance Reduction. *SIAM Journal on Optimization*, 24(4):2057—2075, 2014.
- [30] Yi Xu and Tianbao Yang. First-order Stochastic Algorithms for Escaping From Saddle Points in Almost Linear Time. *ArXiv e-prints*, abs/1711.01944, November 2017.